

# Untitled18

June 26, 2019

## 3. Exploratory Data Analysis

```
In [0]: import warnings
warnings.filterwarnings("ignore")
import shutil
import scipy
import os
import pandas as pd
import matplotlib
matplotlib.use('nbAgg')
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pickle
import xgboost
from sklearn.manifold import TSNE
from sklearn import preprocessing
import pandas as pd
from multiprocessing import Process# this is used for multithreading
import multiprocessing
import codecs# this is used for file operations
import random as r
from xgboost import XGBClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import log_loss
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from tqdm import tqdm
from sklearn.feature_extraction.text import CountVectorizer

In [0]: #separating byte files and asm files

source = 'train'
```

```

destination = 'byteFiles'

# we will check if the folder 'byteFiles' exists if it not there we will create a folder
if not os.path.isdir(destination):
    os.makedirs(destination)

# if we have folder called 'train' (train folder contains both .asm files and .bytes files)
# for every file that we have in our 'asmFiles' directory we check if it is ending with
# 'byteFiles' folder

# so by the end of this snippet we will separate all the .byte files and .asm files
if os.path.isdir(source):
    os.rename(source, 'asmFiles')
    source = 'asmFiles'
    data_files = os.listdir(source)
    for file in data_files:
        if (file.endswith("bytes")):
            shutil.move('asmFiles/'+file, destination)

```

### 3.1. Distribution of malware classes in whole data set


```

In [0]: %matplotlib inline
Y=pd.read_csv("trainLabels.csv")
total = len(Y)*1.
ax=sns.countplot(x="Class", data=Y)
for p in ax.patches:
    ax.annotate('{:.1f}%'.format(100*p.get_height()/total), (p.get_x()+0.1, p.get_height()+0.1))

#put 11 ticks (therefore 10 steps), from 0 to the total number of rows in the dataframe
ax.yaxis.set_ticks(np.linspace(0, total, 11))

#adjust the ticklabel to the desired format, without changing the position of the ticks
ax.set_yticklabels(map('{:.1f}%'.format, 100*ax.yaxis.get_majorticklocs()/total))
plt.show()

```



### 3.2. Feature extraction


In [0]: *#file sizes of byte files*

```
files=os.listdir('byteFiles')
filenames=Y['Id'].tolist()
class_y=Y['Class'].tolist()
class_bytes=[]
sizebytes=[]
fnames=[]
for file in files:
    # print(os.stat('byteFiles/OA32eTdBKayjCWhZqDOQ.txt'))
    # os.stat_result(st_mode=33206, st_ino=1125899906874507, st_dev=3561571700, st_nli
    # st_size=3680109, st_atime=1519638522, st_mtime=1519638522, st_ctime=1519638522)
    # read more about os.stat: here https://www.tutorialspoint.com/python/os_stat.htm
    statinfo=os.stat('byteFiles/'+file)
    # split the file name at '.' and take the first part of it i.e the file name
    file=file.split('.')[0]
    if any(file == filename for filename in filenames):
        i=filenames.index(file)
        class_bytes.append(class_y[i])
        # converting into Mb's
        sizebytes.append(statinfo.st_size/(1024.0*1024.0))
        fnames.append(file)
data_size_byte=pd.DataFrame({'ID':fnames,'size':sizebytes,'Class':class_bytes})
print (data_size_byte.head())
```

	ID	size	Class
0	HJQiyIRqr6FPeBcoaEsk	6.842773	3
1	G8hm6UqIKBQWlMpeTScb	0.801514	1
2	6mUHQtCBjzWA0fGIEnP7	7.596436	2
3	9gMZ6wVFX7KvHN3y8LoG	7.285400	2
4	hqzvHQ4UBkTPinujM1RC	2.308838	6

### 3.2.2 box plots of file size (.byte files) feature

```
In [0]: #boxplot of byte files
ax = sns.boxplot(x="Class", y="size", data=data_size_byte)
plt.title("boxplot of .bytes file sizes")
plt.show()
```



### 3.2.3 feature extraction from byte files

```
In [0]: #removal of addres from byte files
# contents of .byte files
# -----
#00401000 56 8D 44 24 08 50 8B F1 E8 1C 1B 00 00 C7 06 08
#-----
#we remove the starting address 00401000
#removal of addres from byte files
# contents of .byte files
```

```

# -----
#00401000 56 8D 44 24 08 50 8B F1 E8 1C 1B 00 00 C7 06 08
#-----
#we remove the starting address 00401000
import datetime
start = datetime.datetime.now()
files = os.listdir('byteFiles')
filenames=[]
array=[]
for f in files:
    if(f.endswith("bytes")):
        file=f.split('.')[0]
        text_file = open('byteFiles/'+file+".txt", 'w+')
        with open('byteFiles/'+file+".bytes" ,"r") as fp:
            lines=""
            for line in fp:
                a=line.rstrip().split(" ")[1:]
                b=' ' + ' '.join(a)
                b=b+"\n"
                text_file.write(b)
            fp.close()
            os.remove('byteFiles/'+file+".bytes")
        text_file.close()
print("Time required to run this cell:", datetime.datetime.now() - start)

print('done!')

```

```

In [0]: #program to convert into bag of words of bytefiles
#this is custom-built bag of words this is unigram bag of words
byte_feature_file=open('result.csv','w+')
byte_feature_file.write("ID,0,1,2,3,4,5,6,7,8,9,0a,0b,0c,0d,0e,0f,10,11,12,13,14,15,16")
for file in files:
    filenames2.append(f)
    byte_feature_file.write(file+",")
    if(file.endswith("txt")):
        with open('byteFiles/'+file,"r") as byte_flie:
            for lines in byte_flie:
                line=lines.rstrip().split(" ")
                for hex_code in line:
                    if hex_code=='??:
                        feature_matrix[k][256]+=1
                    else:
                        feature_matrix[k][int(hex_code,16)]+=1
            byte_flie.close()
        for i in feature_matrix[k]:
            byte_feature_file.write(str(i)+",")

```

```

byte_feature_file.write("\n")

k += 1

byte_feature_file.close()

In [0]: byte_features=pd.read_csv("result_with_size.csv")
        print (byte_features.head())

```

Unnamed: 0	ID	0	1	2	3	4	5	\
0	0	01azqd4InC7m9JpocGv5	601905	3905	2816	3832	3345	3242
1	1	01IsoiSMh5gxyDYTl4CB	39755	8337	7249	7186	8663	6844
2	2	01jsnpXSAlgW6aPeDxrU	93506	9542	2568	2438	8925	9330
3	3	01kcPWA9K2B0xQeS5Rju	21091	1213	726	817	1257	625
4	4	01SuzwMJEIXsK7A8dQbl	19764	710	302	433	559	410

	6	7	...	f9	fa	fb	fc	fd	fe	ff	??	\
0	3650	3201	...	3101	3211	3097	2758	3099	2759	5753	1824	
1	8420	7589	...	439	281	302	7639	518	17001	54902	8588	
2	9007	2342	...	2242	2885	2863	2471	2786	2680	49144	468	
3	550	523	...	485	462	516	1133	471	761	7998	13940	
4	262	249	...	350	209	239	653	221	242	2199	9008	

```

        size  Class
0  4.234863      9
1  5.538818      2
2  3.887939      9
3  0.574219      1
4  0.370850      8

```

[5 rows x 261 columns]

```

In [0]: data_size_byte.columns

Out[0]: Index(['ID', 'size', 'Class'], dtype='object')

In [0]: result = pd.merge(byte_features, data_size_byte, on='ID', how='left')
        result.head()

```

```

Out[0]:   Unnamed: 0      ID      0      1      2      3      4      5  \
0         0  01azqd4InC7m9JpocGv5  601905  3905  2816  3832  3345  3242
1         1  01IsoiSMh5gxyDYTl4CB   39755  8337  7249  7186  8663  6844
2         2  01jsnpXSAlgW6aPeDxrU   93506  9542  2568  2438  8925  9330
3         3  01kcPWA9K2B0xQeS5Rju   21091  1213   726   817  1257   625
4         4  01SuzwMJEIXsK7A8dQbl   19764   710   302   433   559   410

        6      7      ...      fb      fc      fd      fe      ff      ??      size_x  Class_x  \
0  3650  3201  ...  3097  2758  3099  2759  5753  1824  4.234863      9

```

1	8420	7589	...	302	7639	518	17001	54902	8588	5.538818	2
2	9007	2342	...	2863	2471	2786	2680	49144	468	3.887939	9
3	550	523	...	516	1133	471	761	7998	13940	0.574219	1
4	262	249	...	239	653	221	242	2199	9008	0.370850	8

	size_y	Class_y
0	4.234863	9
1	5.538818	2
2	3.887939	9
3	0.574219	1
4	0.370850	8

[5 rows x 263 columns]

```
In [0]: # result = result.drop(["Class_y", "size_x", "size_y"], axis = 1)
```

```
In [0]: # data = data.rename(columns={"Area": "place_name"})
# result = result.rename(columns={"Class_x": "Class"})
```

```
In [0]: # https://stackoverflow.com/a/29651514
```

```
def normalize(df):
    result1 = df.copy()
    for feature_name in df.columns:
        if (str(feature_name) != str('ID') and str(feature_name) != str('Class')):
            max_value = df[feature_name].max()
            min_value = df[feature_name].min()
            result1[feature_name] = (df[feature_name] - min_value) / (max_value - min_value)
    return result1
result = normalize(result)
```

```
In [0]: result=byte_features
```

```
In [0]: data_y = result['Class']
# split the data into test and train by maintaining same distribution of output variable
X_train, X_test, y_train, y_test = train_test_split(result.drop(['ID', 'Class'], axis=1), data_y, test_size=0.2, random_state=42)
# split the train data into train and cross validation by maintaining same distribution
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, stratify=y_train, test_size=0.2, random_state=42)
```

```
In [0]: dfasm=pd.read_csv("asmoutputfile.csv")
Y.columns = ['ID', 'Class']
result_asm = pd.merge(dfasm, Y,on='ID', how='left')
result_asm.head()
```

```
Out[0]:
```

	ID	HEADER:	.text:	.Pav:	.idata:	.data:	.bss:	\
0	01kcPWA9K2B0xQeS5Rju	19	744	0	127	57	0	
1	1E93CpP6ORHFNiT5Qfvn	17	838	0	103	49	0	
2	3ekVow2ajZHbTnBcsDfX	17	427	0	50	43	0	
3	3X2nY7iQaPBIWDrAZqJe	17	227	0	43	19	0	
4	460ZzdsSKDCFV8h7XWxf	17	402	0	59	170	0	

	.rdata:	.edata:	.rsrc:	...	edx	esi	eax	ebx	ecx	edi	ebp	esp	eip	\
0	323	0	3	...	18	66	15	43	83	0	17	48	29	
1	0	0	3	...	18	29	48	82	12	0	14	0	20	
2	145	0	3	...	13	42	10	67	14	0	11	0	9	
3	0	0	3	...	6	8	14	7	2	0	8	0	6	
4	0	0	3	...	12	9	18	29	5	0	11	0	11	

	Class
0	1
1	1
2	1
3	1
4	1

[5 rows x 53 columns]

In [0]: *#file sizes of byte files*

```
files=os.listdir('asmFiles')
filenames=Y['ID'].tolist()
class_y=Y['Class'].tolist()
class_bytes=[]
sizebytes=[]
fnames=[]
for file in files:
    # print(os.stat('byteFiles/0A32eTdBKayjCWhZqDOQ.txt'))
    # os.stat_result(st_mode=33206, st_ino=1125899906874507, st_dev=3561571700, st_nli
    # st_size=3680109, st_atime=1519638522, st_mtime=1519638522, st_ctime=1519638522)
    # read more about os.stat: here https://www.tutorialspoint.com/python/os_stat.htm
    statinfo=os.stat('asmFiles/'+file)
    # split the file name at '.' and take the first part of it i.e the file name
    file=file.split('.')[0]
    if any(file == filename for filename in filenames):
        i=filenames.index(file)
        class_bytes.append(class_y[i])
        # converting into Mb's
        sizebytes.append(statinfo.st_size/(1024.0*1024.0))
        fnames.append(file)
asm_size_byte=pd.DataFrame({'ID':fnames,'size':sizebytes,'Class':class_bytes})
print (asm_size_byte.head())
```

	ID	size	Class
0	C9hHuINUVJqk1zo5pTQX	8.265899	1
1	16cTMtKIjH5SbyovWuBq	36.532800	2
2	0Phtp2LVcsCFMKkGgmRH	58.857293	9
3	4Z17fDwSI1GLtyoPv6Fp	0.172276	3
4	c8BfFP6iYEIRaUxdGtmX	37.294192	2



```
In [0]: # add the file size feature to previous extracted features
print(result_asm.shape)
print(asm_size_byte.shape)
result_asm = pd.merge(result_asm, asm_size_byte.drop(['Class'], axis=1), on='ID', how='left')
result_asm.head()
```

```
(10868, 53)
```

```
(10868, 3)
```

```
Out[0]:
```

	ID	HEADER:	.text:	.Pav:	.idata:	.data:	.bss:	\
0	01kcPWA9K2B0xQeS5Rju	19	744	0	127	57	0	
1	1E93CpP60RHFNiT5Qfvn	17	838	0	103	49	0	
2	3ekVow2ajZHbTnBcsDfX	17	427	0	50	43	0	
3	3X2nY7iQaPBIWDrAZqJe	17	227	0	43	19	0	
4	460ZzdsSKDCFV8h7XWxf	17	402	0	59	170	0	

	.rdata:	.edata:	.rsrc:	...	esi	eax	ebx	ecx	edi	ebp	esp	eip	\
0	323	0	3	...	66	15	43	83	0	17	48	29	
1	0	0	3	...	29	48	82	12	0	14	0	20	
2	145	0	3	...	42	10	67	14	0	11	0	9	
3	0	0	3	...	8	14	7	2	0	8	0	6	
4	0	0	3	...	9	18	29	5	0	11	0	11	

	Class	size
0	1	0.078190
1	1	0.063400
2	1	0.041695
3	1	0.018757
4	1	0.037567

```
[5 rows x 54 columns]
```

```
In [0]: asm_y = result_asm['Class']
asm_x = result_asm.drop(['ID', 'Class', '.BSS:', 'rtn', '.CODE'], axis=1)
```

```
In [0]: X_train_asm, X_test_asm, y_train_asm, y_test_asm = train_test_split(asm_x, asm_y, stratify=asm_y,
X_train_asm, X_cv_asm, y_train_asm, y_cv_asm = train_test_split(X_train_asm, y_train_asm,
```

```
In [0]: result_x = pd.merge(result, result_asm.drop(['Class'], axis=1), on='ID', how='left')
result_y = result_x['Class']
result_x = result_x.drop(['ID', 'rtn', '.BSS:', '.CODE', 'Class'], axis=1)
result_x.head()
```

```
Out[0]:
```

	Unnamed: 0	0	1	2	3	4	5	6	7	8	...	\
0	0	601905	3905	2816	3832	3345	3242	3650	3201	2965	...	
1	1	39755	8337	7249	7186	8663	6844	8420	7589	9291	...	
2	2	93506	9542	2568	2438	8925	9330	9007	2342	9107	...	
3	3	21091	1213	726	817	1257	625	550	523	1078	...	

```
4          4  19764   710   302   433   559   410   262   249   422   ...
```

```
    edx  esi  eax  ebx  ecx  edi  ebp  esp  eip    size_y
0  808  2290  1281  587  701    0   15   14  456  56.229886
1  260  1090   391  905  420    0   24   22  227  13.999378
2    5   547    5  451   56    0   27    0  117   8.507785
3   18    66   15   43   83    0   17   48   29   0.078190
4   18  1228   24 1546  107    0   15    0   76   0.996723
```

```
[5 rows x 308 columns]
```

```
In [0]: result_y.head()
```

```
Out[0]: 0    9
        1    2
        2    9
        3    1
        4    8
        Name: Class, dtype: int64
```

```
In [0]: X_train, X_test_merge, y_train, y_test_merge = train_test_split(result_x, result_y, stratify=result_y,
                                X_train_merge, X_cv_merge, y_train_merge, y_cv_merge = train_test_split(X_train, y_train,
```

```
In [0]: result_x['ID'] = result.ID
```

## 0.1 Bi-grams

```
In [0]: byte_vocab = "00,01,02,03,04,05,06,07,08,09,0a,0b,0c,0d,0e,0f,10,11,12,13,14,15,16,17,18,19,1a,1b,1c,1d,1e,1f,20,21,22,23,24,25,26,27,28,29,2a,2b,2c,2d,2e,2f,30,31,32,33,34,35,36,37,38,39,3a,3b,3c,3d,3e,3f,40,41,42,43,44,45,46,47,48,49,4a,4b,4c,4d,4e,4f,50,51,52,53,54,55,56,57,58,59,5a,5b,5c,5d,5e,5f,60,61,62,63,64,65,66,67,68,69,6a,6b,6c,6d,6e,6f,70,71,72,73,74,75,76,77,78,79,7a,7b,7c,7d,7e,7f,80,81,82,83,84,85,86,87,88,89,8a,8b,8c,8d,8e,8f,90,91,92,93,94,95,96,97,98,99,a0,a1,a2,a3,a4,a5,a6,a7,a8,a9,aa,ab,ac,ad,ae,af,b0,b1,b2,b3,b4,b5,b6,b7,b8,b9,ba,bb,bc,bd,be,bf,c0,c1,c2,c3,c4,c5,c6,c7,c8,c9,ca,cb,cc,cd,ce,cf,d0,d1,d2,d3,d4,d5,d6,d7,d8,d9,da,db,dc,dd,de,df,e0,e1,e2,e3,e4,e5,e6,e7,e8,e9,ea,eb,ec,ed,ee,ef,f0,f1,f2,f3,f4,f5,f6,f7,f8,f9,fa,fb,fc,fd,fe,ff"
```

```
In [0]: byte_bigram_vocab = []
```

```
def byte_bigram():
    for i, v in enumerate(byte_vocab.split(',')):
        for j in range(0, len(byte_vocab.split(','))):
            byte_bigram_vocab.append(v + ' ' + byte_vocab.split(',')[j])
    len(byte_bigram_vocab)
```

```
In [0]: byte_bigram()
```

```
In [0]: byte_bigram_vocab[:5]
```

```
Out[0]: ['00 00', '00 01', '00 02', '00 03', '00 04']
```

```
In [0]: len(byte_bigram_vocab)
```

```
Out[0]: 66049
```

```
In [0]: byte_trigram_vocab = []
```

```
def byte_trigram():
```

```

        for i, v in enumerate(byte_vocab.split(',')):
            for j in range(0, len(byte_vocab.split(','))):
                for k in range(0, len(byte_vocab.split(','))):
                    byte_trigram_vocab.append(v + ' ' + byte_vocab.split(',')[j] + ' ' + byte_vocab.split(',')[k])
        len(byte_trigram_vocab)

In [0]: byte_trigram()

In [0]: import pickle

        filename = 'trigram'
        outfile = open(filename, 'wb')

        pickle.dump(byte_trigram_vocab, outfile)
        outfile.close()

In [0]: infile = open('trigram', 'rb')
        byte_trigram_vocab = pickle.load(infile)
        infile.close()

In [0]: byte_trigram_vocab[:5]

Out[0]: ['00 00 00', '00 00 01', '00 00 02', '00 00 03', '00 00 04']

In [0]: len(byte_trigram_vocab)

Out[0]: 16974593

In [0]: from tqdm import tqdm
        from sklearn.feature_extraction.text import CountVectorizer

        vect = CountVectorizer(lowercase=False, ngram_range=(2,2), vocabulary=byte_bigram_vocab)
        byte_bigram_vect = scipy.sparse.csr_matrix((10868, 66049))
        for i, file in tqdm(enumerate(os.listdir('./byteFiles'))):
            f = open('./byteFiles/' + file)
            a[i : ] += scipy.sparse.csr_matrix(vect.fit_transform([f.read().replace('\n', ' ')]))
            f.close()

        scipy.sparse.save_npz('bytebigram.npz', byte_bigram_vect)

0it [00:00, ?it/s]

1it [00:01, 1.01s/it]

2it [00:01, 1.28it/s]

```

```

3it [00:03, 1.15s/it]
4it [00:05, 1.40s/it]
5it [00:05, 1.19s/it]
6it [00:06, 1.11s/it]
7it [00:06, 1.23it/s]
8it [00:08, 1.02it/s]
9it [00:08, 1.38it/s]
10it [00:10, 1.10s/it]
11it [00:12, 1.27s/it]
12it [00:14, 1.45s/it]
13it [00:15, 1.51s/it]
14it [00:16, 1.22s/it]
15it [00:16, 1.09it/s]
16it [00:17, 1.09it/s]
17it [00:17, 1.38it/s]
18it [00:17, 1.70it/s]
19it [00:19, 1.30it/s]
20it [00:20, 1.20it/s]
21it [00:20, 1.61it/s]

```

```
In [0]: scipy.sparse.save_npz('bytebigram.npz', byte_bigram_vect)
```

```
In [0]: import scipy
        from sklearn.preprocessing import normalize
        byte_bigram_vect = normalize(scipy.sparse.load_npz('bytebigram.npz'), axis = 0)
```

# 1 N-Grams

```
In [0]: #Ref https://www.edwardraff.com/publications/what\_can\_ngrams\_learn.pdf  
#Ref https://github.com/melanieihuei/Malware-Classification
```

```
opcodes = ['jmp', 'mov', 'retf', 'push', 'pop', 'xor', 'retn', 'nop', 'sub', 'inc', 'd
```

```
In [0]: asm_bigram = []  
def asmopcodebigram():  
    for i, v in enumerate(opcodes):  
        for j in range(0, len(opcodes)):  
            asm_bigram.append(v + ' ' + opcodes[j])
```

```
In [0]: asmopcodebigram()  
len(asm_bigram)
```

```
Out[0]: 676
```

```
In [0]: asm_trigram = []  
def asmopcodetrigram():  
    for i, v in enumerate(opcodes):  
        for j in range(0, len(opcodes)):  
            for k in range(0, len(opcodes)):  
                asm_trigram.append(v + ' ' + opcodes[j] + ' ' + opcodes[k])
```

```
In [0]: asmopcodetrigram()  
len(asm_trigram)
```

```
Out[0]: 17576
```

```
In [0]: asm_4gram = []  
for i, v in enumerate(opcodes):  
    for j in range(0, len(opcodes)):  
        for k in range(0, len(opcodes)):  
            for l in range(0, len(opcodes)):  
                asm_4gram.append(v + ' ' + opcodes[j] + ' ' + opcodes[k] + ' ' + opcodes[l])  
len(asm_4gram)
```

```
Out[0]: 456976
```

```
In [0]: def opcode_collect():  
    op_file = open("opcode_file.txt", "w+")  
    for asmfile in os.listdir('asmFiles'):  
        opcode_str = ""  
        with codecs.open('asmFiles/' + asmfile, encoding='cp1252', errors='replace') as f:  
            for lines in f:  
                line = lines.rstrip().split()  
                for li in line:  
                    if li in opcodes:
```

```

        opcode_str += li + ' '
        op_file.write(opcode_str + "\n")
    op_file.close()
opcode_collect()

In [0]: vect = CountVectorizer(ngram_range=(2, 2), vocabulary = asm_bigram)
bigram_vect = scipy.sparse.csr_matrix((10868, len(asm_bigram)))
raw_opcode = open('opcode_file.txt').read().split('\n')

for i in range(10868):
    bigram_vect[i, :] += scipy.sparse.csr_matrix(vect.transform([raw_opcode[i]]))

In [0]: scipy.sparse.save_npz('op_bigram.npz', opcodebivect)

In [0]: vect = CountVectorizer(ngram_range=(3, 3), vocabulary = asm_trigram)
trigram_vect = scipy.sparse.csr_matrix((10868, len(asm_trigram)))
raw_opcode = open('opcode_file.txt').read().split('\n')


for i in range(10868):
    trigram_vect[i, :] += scipy.sparse.csr_matrix(vect.transform([raw_opcode[i]]))

In [0]: scipy.sparse.save_npz('op_trigram.npz', opcodetrivect)

In [0]: opcodebivect=scipy.sparse.load_npz('op_bigram.npz')
opcodetrivect=scipy.sparse.load_npz('op_trigram.npz')

In [0]: bi_imp_feat = imp_features(normalize(bigram_vect, axis = 0), asm_bigram, 200)


```



```
In [0]: op_bigram = pd.read_csv('op_bigram.csv').drop('Unnamed: 0', axis = 1).fillna(0)
        op_trigram = pd.read_csv('op_trigram.csv').drop('Unnamed: 0', axis = 1).fillna(0)

In [0]: opcode_tetra_vect = scipy.sparse.load_npz('opcode_tetragram.npz')

In [0]: n_imp_feat = imp_features(normalize(n_vect, axis = 0), asm_4gram, 5000)
```



```
In [0]: op_tetragram = pd.SparseDataFrame(normalize(n_vect, axis = 0), columns = asm_4gram)
        op_tetragram = op_tetragram.loc[:, np.intersect1d(op_tetragram.columns, np.take(asm_4g

In [0]: op_tetra_df.to_dense().to_csv('op_tetra_filtered.csv')

In [0]: op_tetragram = pd.read_csv('op_n_final.csv').drop('Unnamed: 0', axis = 1).fillna(0)

In [0]: op_tetragram['ID'] = result.ID
        op_tetragram.head()
```

```
Out[0]:   add add add add   add add add cmp   add add add dec   add add add jmp \
0         0.000443         0.024936         0.0         0.0
1         0.000000         0.000000         0.0         0.0
2         0.000000         0.000000         0.0         0.0
3         0.000000         0.000000         0.0         0.0
4         0.000000         0.000000         0.0         0.0

        add add add jz   add add add mov   add add add or   add add add pop \
```

0	0.0	0.003054	0.0	0.0
1	0.0	0.000000	0.0	0.0
2	0.0	0.000000	0.0	0.0
3	0.0	0.000000	0.0	0.0
4	0.0	0.000000	0.0	0.0

	add	add	add	push	add	add	add	retn	...	xor	xor	push	push	\
0				0.008047				0.0	...				0.0	
1				0.000000				0.0	...				0.0	
2				0.000000				0.0	...				0.0	
3				0.000000				0.0	...				0.0	
4				0.000000				0.0	...				0.0	

	xor	xor	push	sub	xor	xor	push	xor	xor	xor	sub	mov	xor	xor	sub	push	\
0				0.0				0.0				0.0				0.0	
1				0.0				0.0				0.0				0.0	
2				0.0				0.0				0.0				0.0	
3				0.0				0.0				0.0				0.0	
4				0.0				0.0				0.0				0.0	

	xor	xor	xchg	mov	xor	xor	xor	mov	xor	xor	xor	sub	xor	xor	xor	xor	\
0				0.0				0.000216				0.0				0.00111	
1				0.0				0.000000				0.0				0.00000	
2				0.0				0.000000				0.0				0.00000	
3				0.0				0.000000				0.0				0.00000	
4				0.0				0.000000				0.0				0.00000	

	ID
0	01azqd4InC7m9JpocGv5
1	01IsoiSMh5gxyDYTl4CB
2	01jsnpXSAlgw6aPeDxrU
3	01kcPWA9K2B0xQeS5Rju
4	01SuzwMJEIXsK7A8dQbl

[5 rows x 5001 columns]


```
In [0]: def imp_features(data, features, keep):
    rf = RandomForestClassifier(n_estimators = 100, n_jobs = -1)
    rf.fit(data, result_y)
    imp_feature_indx = np.argsort(rf.feature_importances_)[::-1]
    imp_value = np.take(rf.feature_importances_, imp_feature_indx[:20])
    imp_feature_name = np.take(features, imp_feature_indx[:20])
    sns.set()
    plt.figure(figsize = (10, 5))
    ax = sns.barplot(x = imp_feature_name, y = imp_value)
    ax.set_xticklabels(labels = imp_feature_name, rotation = 45)
    sns.set_palette(reversed(sns.color_palette("husl", 10)), 10)
    plt.title('Important Features')
```



```
plt.xlabel('Feature Names')
plt.ylabel('Importance')
return imp_feature_indx[:keep]
```

```
In [0]: from sklearn.preprocessing import normalize
```

```
tri_imp_feat = imp_features(normalize(bigram_vect, axis = 0), asm_trigram, 300)
```



```
In [0]: op_bigram = pd.SparseDataFrame(normalize(bigram_vect, axis = 0), columns = asm_bigram)
for col in op_bigram.columns:
    if col not in np.take(asm_bigram, bi_imp_feat):
        op_bigram.drop(col, axis = 1, inplace = True)
```

```
In [0]: op_bigram = pd.read_csv('op_bigram_final.csv').drop('Unnamed: 0', axis = 1).fillna(0)
op_bigram['ID'] = result.ID
op_bigram.head()
```

```
Out[0]:
```

	jmp jmp	jmp mov	jmp push	jmp pop	jmp xor	jmp sub	jmp add	\
0	0.002169	0.016612	0.015480	0.002994	0.017598	0.025732	0.022664	
1	0.009038	0.001400	0.002101	0.000374	0.006425	0.000000	0.000000	
2	0.031815	0.003894	0.000420	0.000000	0.002374	0.008950	0.016752	
3	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
4	0.009255	0.001095	0.002101	0.000374	0.005587	0.000000	0.000000	

	jmp cmp	jmp call	jmp lea	...	lea add	lea or	lea cmp	lea call	\
0	0.022002	0.013257	0.005995	...	0.013914	0.006827	0.017680	0.003501	

1	0.002806	0.001128	0.000521	...	0.001113	0.002276	0.003584	0.000400
2	0.000112	0.000564	0.000391	...	0.002226	0.007396	0.001911	0.000833
3	0.000000	0.000000	0.000000	...	0.001113	0.002276	0.000478	0.000000
4	0.000449	0.001410	0.000521	...	0.000000	0.000000	0.001911	0.000333

	lea jz	lea lea	movzx mov	movzx sub	movzx add	ID
0	0.010045	0.007178	0.015875	0.056518	0.004574	01azqd4InC7m9JpocGv5
1	0.000000	0.000479	0.000772	0.036972	0.000286	01IsoiSMh5gxyDYT14CB
2	0.002318	0.006102	0.000000	0.000000	0.000000	01jsnpXSAlgW6aPeDxrU
3	0.000000	0.000000	0.000000	0.000000	0.000000	01kcPWA9K2B0xQeS5Rju
4	0.000000	0.000120	0.000000	0.036266	0.000000	01SuzwMJEIXsK7A8dQb1

[5 rows x 201 columns]

```
In [0]: op_trigram = pd.SparseDataFrame(normalize(trigram_vect, axis = 0), columns = asm_trigram)
op_trigram = op_trigram.loc[:, np.intersect1d(op_trigram.columns, np.take(asm_trigram,
```

```
In [0]: op_trigram = pd.read_csv('op_trigram_final.csv').drop('Unnamed: 0', axis = 1).fillna(0)
op_trigram['ID'] = result.ID
op_trigram.head()
```

```
Out[0]:      jmp add add      jmp add call      jmp add cmp      jmp add dec      jmp add imul  \
0          0.0          0.033763          0.026559          0.0          0.0
1          0.0          0.000000          0.000000          0.0          0.0
2          0.0          0.000000          0.013280          0.0          0.0
3          0.0          0.000000          0.000000          0.0          0.0
4          0.0          0.000000          0.000000          0.0          0.0
```

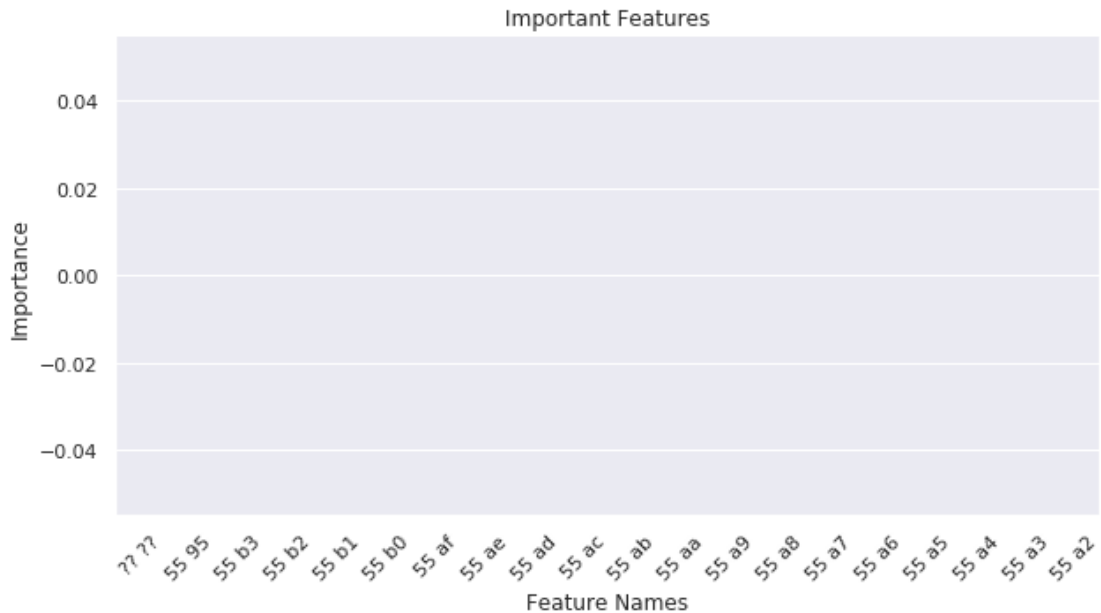
	jmp add inc	jmp add jmp	jmp add jz	jmp add lea	jmp add mov	...	\
0	0.042737	0.000000	0.013761	0.033191	0.023325	...	
1	0.000000	0.000000	0.000000	0.000000	0.000000	...	
2	0.000000	0.005556	0.000000	0.000000	0.026069	...	
3	0.000000	0.000000	0.000000	0.000000	0.000000	...	
4	0.000000	0.000000	0.000000	0.000000	0.000000	...	

	jmp xor lea	jmp xor mov	jmp xor or	jmp xor pop	jmp xor push	\
0	0.006909	0.007039	0.018413	0.010029	0.011196	
1	0.006909	0.011856	0.000000	0.002006	0.000000	
2	0.000000	0.001852	0.000000	0.000000	0.001120	
3	0.000000	0.000000	0.000000	0.000000	0.000000	
4	0.000000	0.011856	0.000000	0.002006	0.000000	

	jmp xor ret	jmp xor sub	jmp xor xchg	jmp xor xor	ID
0	0.000000	0.0	0.0	0.008479	01azqd4InC7m9JpocGv5
1	0.003717	0.0	0.0	0.012718	01IsoiSMh5gxyDYT14CB
2	0.003717	0.0	0.0	0.004239	01jsnpXSAlgW6aPeDxrU
3	0.000000	0.0	0.0	0.000000	01kcPWA9K2B0xQeS5Rju
4	0.003717	0.0	0.0	0.004239	01SuzwMJEIXsK7A8dQb1

[5 rows x 301 columns]

```
In [0]: byte_imp_feat = imp_features(normalize(byte_vect, axis = 0), byte_bigram_vocab, 300)
```



```
In [0]: byte_imp_feat = np.load('byte_imp_feat.npy')
```

```
byte_imp = np.zeros((10868, 0))
for i in byte_imp:
    sliced = byte_vect[:, i].todense()
    byte_imp = np.hstack([byte_imp, sliced])
```

```
In [0]: byte_imp = pd.SparseDataFrame(byte_imp, columns = np.take(byte_bigram_vocab, byte_imp_
```

```
In [0]: byte_bigram = pd.read_csv('byte_bigram.csv').drop('Unnamed: 0', axis = 1).fillna(0)
```

```
byte_bigram['ID'] = result.ID
```

```
byte_bigram.head()
```

```
Out[0]:
```

	?? ??	55 95	55 b3	55 b2	55 b1	55 b0	55 af	55 ae	55 ad	55 ac	...	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	

	54 b3	54 b4	54 c4	54 d1	54 d0	54 cf	54 ce	54 cd	54 cc	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	ID
0	01azqd4InC7m9JpocGv5
1	01IsoiSMh5gxyDYTl4CB
2	01jsnpXSAlgW6aPeDxrU
3	01kcPWA9K2B0xQeS5Rju
4	01SuzwMJEIXsK7A8dQbl

[5 rows x 301 columns]

## 1.1 Top 200 Image Features


```
In [0]: #Ref https://github.com/adeya99/Microsoft-Malware-Detection/blob/master/Malware%20Classification%20Features%20Top%20200%20Features.txt
import array
```

```
def collect_img_asm():
    #pix_file = open("../pixels.txt", "w+")
    for asmfile in os.listdir("./asmFiles"):
        file_name = asmfile.split('.')[0]
        file = codecs.open("./asmFiles/" + asmfile, 'rb')
        file_len = os.path.getsize("./asmFiles/" + asmfile)
        width = int(file_len ** 0.5)
        rem = int(file_len / width)
        arr = array.array('B')
        arr.frombytes(file.read())
        file.close()
        reshaped = np.reshape(arr[:width * width], (width, width))
        reshaped = np.uint8(reshaped)
        scipy.misc.imsave('./asm_image/' + file_name + '.png', reshaped)
```

```
collect_img_asm()
```

```
In [0]: from IPython.display import Image
Image(filename='asm_image/8vJiQRcq15ZtmEdHOIp.png')
```

Out[0]:



```
In [0]: import cv2
        image_features = np.zeros((10868, 200))

        for i, asmfile in enumerate(os.listdir("asmFiles")):
            img = cv2.imread("asm_image/" + asmfile.split('.')[0] + '.png')
            img_arr = img.flatten()[:200]
            image_features[i, :] += img_arr

In [0]: from sklearn.preprocessing import normalize

        img_feat = []
        for i in range(200):
            img_feat.append('pix' + str(i))
        img_final = pd.DataFrame(normalize(image_features, axis = 0), columns = img_feat)
```

```
In [0]: img_final['ID'] = result.ID
img_final.head()
```

```
Out [0]:
```

	pix0	pix1	pix2	pix3	pix4	pix5	pix6	\
0	0.010268	0.010268	0.010268	0.008033	0.008033	0.008033	0.008320	
1	0.006560	0.006560	0.006560	0.013504	0.013504	0.013504	0.012927	
2	0.010268	0.010268	0.010268	0.008033	0.008033	0.008033	0.008320	
3	0.010268	0.010268	0.010268	0.008033	0.008033	0.008033	0.008320	
4	0.006560	0.006560	0.006560	0.013504	0.013504	0.013504	0.012927	

	pix7	pix8	pix9	...	pix191	pix192	pix193	pix194	\
0	0.008320	0.008320	0.007913	...	0.009593	0.009593	0.009593	0.009593	
1	0.012927	0.012927	0.013963	...	0.009593	0.009593	0.009593	0.009593	
2	0.008320	0.008320	0.007913	...	0.009593	0.009593	0.009593	0.009593	
3	0.008320	0.008320	0.007913	...	0.009593	0.009593	0.009593	0.009593	
4	0.012927	0.012927	0.013963	...	0.009593	0.009593	0.009593	0.009593	

	pix195	pix196	pix197	pix198	pix199	ID
0	0.009593	0.009593	0.009593	0.009593	0.009593	01azqd4InC7m9JpocGv5
1	0.009593	0.009593	0.009593	0.009593	0.009593	01IsoiSMh5gxyDYTl4CB
2	0.009593	0.009593	0.009593	0.009593	0.009593	01jsnpXSAlgW6aPeDxrU
3	0.009593	0.009593	0.009593	0.009593	0.009593	01kcPWA9K2B0xQeS5Rju
4	0.009593	0.009593	0.009593	0.009593	0.009593	01SuzwMJEIXsK7A8dQbl

[5 rows x 201 columns]

```
In [0]: img_final.head()
```

```
Out [0]:
```

	pix0	pix1	pix2	pix3	pix4	pix5	pix6	\
0	0.010268	0.010268	0.010268	0.008033	0.008033	0.008033	0.008320	
1	0.006560	0.006560	0.006560	0.013504	0.013504	0.013504	0.012927	
2	0.010268	0.010268	0.010268	0.008033	0.008033	0.008033	0.008320	
3	0.010268	0.010268	0.010268	0.008033	0.008033	0.008033	0.008320	
4	0.006560	0.006560	0.006560	0.013504	0.013504	0.013504	0.012927	

	pix7	pix8	pix9	...	pix191	pix192	pix193	pix194	\
0	0.008320	0.008320	0.007913	...	0.009593	0.009593	0.009593	0.009593	
1	0.012927	0.012927	0.013963	...	0.009593	0.009593	0.009593	0.009593	
2	0.008320	0.008320	0.007913	...	0.009593	0.009593	0.009593	0.009593	
3	0.008320	0.008320	0.007913	...	0.009593	0.009593	0.009593	0.009593	
4	0.012927	0.012927	0.013963	...	0.009593	0.009593	0.009593	0.009593	

	pix195	pix196	pix197	pix198	pix199	ID
0	0.009593	0.009593	0.009593	0.009593	0.009593	01azqd4InC7m9JpocGv5
1	0.009593	0.009593	0.009593	0.009593	0.009593	01IsoiSMh5gxyDYTl4CB
2	0.009593	0.009593	0.009593	0.009593	0.009593	01jsnpXSAlgW6aPeDxrU
3	0.009593	0.009593	0.009593	0.009593	0.009593	01kcPWA9K2B0xQeS5Rju
4	0.009593	0.009593	0.009593	0.009593	0.009593	01SuzwMJEIXsK7A8dQbl

[5 rows x 201 columns]

```
In [0]: # final_data = pd.concat([result_x, op_bi_df, op_tri_df, byte_bi_df, img_df], axis = 1)
final_data = pd.concat([result_x, op_bigram, op_trigram, op_tetragram, byte_bigram, img
```

```
In [0]: final_data = final_data.drop('ID', axis = 1)
final_data.head()
```

```
Out[0]:
```

	Unnamed: 0	0	1	2	3	4	5	6	7	8	...	\
0	0	601905	3905	2816	3832	3345	3242	3650	3201	2965	...	
1	1	39755	8337	7249	7186	8663	6844	8420	7589	9291	...	
2	2	93506	9542	2568	2438	8925	9330	9007	2342	9107	...	
3	3	21091	1213	726	817	1257	625	550	523	1078	...	
4	4	19764	710	302	433	559	410	262	249	422	...	

	pix190	pix191	pix192	pix193	pix194	pix195	pix196	\
0	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	
1	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	
2	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	
3	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	
4	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	0.009593	

	pix197	pix198	pix199
0	0.009593	0.009593	0.009593
1	0.009593	0.009593	0.009593
2	0.009593	0.009593	0.009593
3	0.009593	0.009593	0.009593
4	0.009593	0.009593	0.009593

[5 rows x 6308 columns]

```
In [0]: # final_data = pd.read_csv('final_data_img.csv')
final_data = pd.read_csv('final_data_img4.csv')
```

```
In [0]: X_train_final, X_test_final, y_train_final, y_test_final = train_test_split(final_data,
X_trn_final, X_cv_final, y_trn_final, y_cv_final = train_test_split(X_train_final, y_t
```

#### 4. Machine Learning Models

##### Without 2,3 or 4 gram

```
In [0]: # final_data = pd.read_csv('final_data_img.csv')
final_data = pd.read_csv('final_data.csv')
```

```
In [0]: X_train_final, X_test_final, y_train_final, y_test_final = train_test_split(final_data,
X_trn_final, X_cv_final, y_trn_final, y_cv_final = train_test_split(X_train_final, y_t
```

##### 4.1.1. Random Model

```

In [0]: # we need to generate 9 numbers and the sum of numbers should be 1
        # one solution is to generate 9 numbers and divide each of the numbers by their sum
        # ref: https://stackoverflow.com/a/18662466/4084039

        test_data_len = X_test.shape[0]
        cv_data_len = X_cv.shape[0]

        # we create a output array that has exactly same size as the CV data
        cv_predicted_y = np.zeros((cv_data_len,9))
        for i in range(cv_data_len):
            rand_probs = np.random.rand(1,9)
            cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
        print("Log loss on Cross Validation Data using Random Model",log_loss(y_cv,cv_predicted_y))

        # Test-Set error.
        #we create a output array that has exactly same as the test data
        test_predicted_y = np.zeros((test_data_len,9))
        for i in range(test_data_len):
            rand_probs = np.random.rand(1,9)
            test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
        print("Log loss on Test Data using Random Model",log_loss(y_test,test_predicted_y, epsilon=1e-6))

        predicted_y = np.argmax(test_predicted_y, axis=1)
        plot_confusion_matrix(y_test, predicted_y+1)

```

Log loss on Cross Validation Data using Random Model 2.45615644965

Log loss on Test Data using Random Model 2.48503905509

Number of misclassified points 88.5004599816

----- Confusion matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

----- Precision matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Sum of columns in precision matrix [ 1. 1. 1. 1. 1. 1. 1. 1. 1.]

----- Recall matrix -----



<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Sum of rows in precision matrix [ 1. 1. 1. 1. 1. 1. 1. 1. 1.]

#### 4.1.2. K Nearest Neighbour Classification

```
In [0]: # find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/g
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights=uniform, algorithm=auto, leaf_size=30, p
# metric=minkowski, metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.applidaicourse.com/course/applied-ai-course-online/lessons/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/module
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=3,
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model
# get_params([deep])                  Get parameters for this estimator.
# predict(X)                          Predict the target of new samples.
# predict_proba(X)                    Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [x for x in range(1, 15, 2)]
cv_log_error_array=[]
for i in alpha:
    k_cfl=KNeighborsClassifier(n_neighbors=i)
    k_cfl.fit(X_train,y_train)
    sig_clf = CalibratedClassifierCV(k_cfl, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_cv)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=k_cfl.classes_, eps=1e-
```

```

for i in range(len(cv_log_error_array)):
    print ('log_loss for k = ',alpha[i],'is',cv_log_error_array[i])

best_alpha = np.argmin(cv_log_error_array)

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

k_cfl=KNeighborsClassifier(n_neighbors=alpha[best_alpha])
k_cfl.fit(X_train,y_train)
sig_clf = CalibratedClassifierCV(k_cfl, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print ('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(X_train,predict_y))
predict_y = sig_clf.predict_proba(X_cv)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss(X_cv,predict_y))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(X_test,predict_y))
plot_confusion_matrix(y_test, sig_clf.predict(X_test))

```

```

log_loss for k = 1 is 0.225386237304
log_loss for k = 3 is 0.230795229168
log_loss for k = 5 is 0.252421408646
log_loss for k = 7 is 0.273827486888
log_loss for k = 9 is 0.286469181555
log_loss for k = 11 is 0.29623391147
log_loss for k = 13 is 0.307551203154

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

For values of best alpha = 1 The train log loss is: 0.0782947669247
For values of best alpha = 1 The cross validation log loss is: 0.225386237304
For values of best alpha = 1 The test log loss is: 0.241508604195
Number of misclassified points 4.50781968721

```

----- Confusion matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

----- Precision matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Sum of columns in precision matrix [ 1. 1. 1. 1. 1. 1. 1. 1. 1.]

----- Recall matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Sum of rows in precision matrix [ 1. 1. 1. 1. 1. 1. 1. 1. 1.]

### 4.1.3. Logistic Regression

```
In [0]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/
#-----

alpha = [10 ** x for x in range(-5, 4)]
cv_log_error_array=[]
```

```

for i in alpha:
    logisticR=LogisticRegression(penalty='l2',C=i,class_weight='balanced')
    logisticR.fit(X_train,y_train)
    sig_clf = CalibratedClassifierCV(logisticR, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_cv)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=logisticR.classes_, eps=1e-15))

for i in range(len(cv_log_error_array)):
    print ('log_loss for c = ',alpha[i],'is',cv_log_error_array[i])

best_alpha = np.argmin(cv_log_error_array)

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

logisticR=LogisticRegression(penalty='l2',C=alpha[best_alpha],class_weight='balanced')
logisticR.fit(X_train,y_train)
sig_clf = CalibratedClassifierCV(logisticR, method="sigmoid")
sig_clf.fit(X_train, y_train)
pred_y=sig_clf.predict(X_test)

predict_y = sig_clf.predict_proba(X_train)
print ('log loss for train data',log_loss(y_train, predict_y, labels=logisticR.classes_))
predict_y = sig_clf.predict_proba(X_cv)
print ('log loss for cv data',log_loss(y_cv, predict_y, labels=logisticR.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(X_test)
print ('log loss for test data',log_loss(y_test, predict_y, labels=logisticR.classes_, eps=1e-15))
plot_confusion_matrix(y_test, sig_clf.predict(X_test))

log_loss for c = 1e-05 is 1.56916911178
log_loss for c = 0.0001 is 1.57336384417
log_loss for c = 0.001 is 1.53598598273
log_loss for c = 0.01 is 1.01720972418
log_loss for c = 0.1 is 0.857766083873
log_loss for c = 1 is 0.711154393309
log_loss for c = 10 is 0.583929522635
log_loss for c = 100 is 0.549929846589
log_loss for c = 1000 is 0.624746769121

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
log loss for train data 0.498923428696
log loss for cv data 0.549929846589
log loss for test data 0.528347316704
Number of misclassified points 12.3275068997
```

----- Confusion matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

----- Precision matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
Sum of columns in precision matrix [ 1.  1.  1.  1. nan  1.  1.  1.  1.]
```

----- Recall matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
Sum of rows in precision matrix [ 1.  1.  1.  1.  1.  1.  1.  1.  1.]
```

#### 4.1.4. Random Forest Classifier

```
In [0]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion=gini, max_depth=N
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=auto, max_leaf_nodes=
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None
# class_weight=None)

# Some of methods of RandomForestClassifier()
```

```

# fit(X, y, [sample_weight])          Fit the SVM model according to the given training
# predict(X)                          Perform classification on samples in X.
# predict_proba (X)                  Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/
# -----

alpha=[10,50,100,500,1000,2000,3000]
cv_log_error_array=[]
train_log_error_array=[]
from sklearn.ensemble import RandomForestClassifier
for i in alpha:
    r_cfl=RandomForestClassifier(n_estimators=i,random_state=42,n_jobs=-1)
    r_cfl.fit(X_train,y_train)
    sig_clf = CalibratedClassifierCV(r_cfl, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_cv)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=r_cfl.classes_, eps=1e-10))

for i in range(len(cv_log_error_array)):
    print ('log_loss for c = ',alpha[i],'is',cv_log_error_array[i])

best_alpha = np.argmin(cv_log_error_array)

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

r_cfl=RandomForestClassifier(n_estimators=alpha[best_alpha],random_state=42,n_jobs=-1)
r_cfl.fit(X_train,y_train)
sig_clf = CalibratedClassifierCV(r_cfl, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(y_train, predict_y, labels=r_cfl.classes_, eps=1e-10))

```

```

predict_y = sig_clf.predict_proba(X_cv)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss")
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y))
plot_confusion_matrix(y_test, sig_clf.predict(X_test))

log_loss for c = 10 is 0.106357709164
log_loss for c = 50 is 0.0902124124145
log_loss for c = 100 is 0.0895043339776
log_loss for c = 500 is 0.0881420869288
log_loss for c = 1000 is 0.0879849524621
log_loss for c = 2000 is 0.0881566647295
log_loss for c = 3000 is 0.0881318948443

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

For values of best alpha = 1000 The train log loss is: 0.0266476291801
For values of best alpha = 1000 The cross validation log loss is: 0.0879849524621
For values of best alpha = 1000 The test log loss is: 0.0858346961407
Number of misclassified points 2.02391904324
----- Confusion matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

----- Precision matrix -----

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Sum of columns in precision matrix [ 1.  1.  1.  1.  1.  1.  1.  1.  1.]
----- Recall matrix -----

<IPython.core.display.Javascript object>

```

<IPython.core.display.HTML object>

Sum of rows in precision matrix [ 1. 1. 1. 1. 1. 1. 1. 1. 1.]

#### 4.1.5. XgBoost Classification

#### 4.5.4. Random Forest Classifier on final features

```
In [0]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion=gini, max_depth=N
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=auto, max_leaf_nodes=
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])          Fit the SVM model according to the given training
# predict(X)                          Perform classification on samples in X.
# predict_proba (X)                  Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/
# -----

alpha=[10,50,100,500,1000,2000,3000]
cv_log_error_array=[]
from sklearn.ensemble import RandomForestClassifier
for i in alpha:
    r_cfl=RandomForestClassifier(n_estimators=i,random_state=42,n_jobs=-1)
    r_cfl.fit(X_train_merge,y_train_merge)
    sig_clf = CalibratedClassifierCV(r_cfl, method="sigmoid")
    sig_clf.fit(X_train_merge, y_train_merge)
    predict_y = sig_clf.predict_proba(X_cv_merge)
    cv_log_error_array.append(log_loss(y_cv_merge, predict_y, labels=r_cfl.classes_, e

for i in range(len(cv_log_error_array)):
    print ('log_loss for c = ',alpha[i], 'is',cv_log_error_array[i])

best_alpha = np.argmin(cv_log_error_array)

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
```



```

for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

r_cfl=RandomForestClassifier(n_estimators=alpha[best_alpha],random_state=42,n_jobs=-1)
r_cfl.fit(X_train_merge,y_train_merge)
sig_clf = CalibratedClassifierCV(r_cfl, method="sigmoid")
sig_clf.fit(X_train_merge, y_train_merge)

predict_y = sig_clf.predict_proba(X_train_merge)
print ('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(y_train_merge,predict_y))
predict_y = sig_clf.predict_proba(X_cv_merge)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss(y_cv_merge,predict_y))
predict_y = sig_clf.predict_proba(X_test_merge)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(y_test_merge,predict_y))

log_loss for c = 10 is 0.0461221662017
log_loss for c = 50 is 0.0375229563452
log_loss for c = 100 is 0.0359765822455
log_loss for c = 500 is 0.0358291883873
log_loss for c = 1000 is 0.0358403093496
log_loss for c = 2000 is 0.0357908022178
log_loss for c = 3000 is 0.0355909487962

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

For values of best alpha = 3000 The train log loss is: 0.0166267614753
For values of best alpha = 3000 The cross validation log loss is: 0.0355909487962
For values of best alpha = 3000 The test log loss is: 0.0401141303589

```

#### 4.5.5. XgBoost Classifier on final features

In [0]: # Training a hyper-parameter tuned Xg-Boost regressor on our train data

```

# find more about XGBClassifier function here http://xgboost.readthedocs.io/en/latest/
# -----
# default paramters
# class xgboost.XGBClassifier(max_depth=3, learning_rate=0.1, n_estimators=100, silent=0)

```

```

# objective='binary:logistic', booster='gbtree', n_jobs=1, nthread=None, gamma=0, min_
# max_delta_step=0, subsample=1, colsample_bytree=1, colsample_bylevel=1, reg_alpha=0,
# scale_pos_weight=1, base_score=0.5, random_state=0, seed=None, missing=None, **kwargs

# some of methods of RandomForestRegressor()
# fit(X, y, sample_weight=None, eval_set=None, eval_metric=None, early_stopping_rounds=
# get_params([deep])          Get parameters for this estimator.
# predict(data, output_margin=False, ntree_limit=0) : Predict with data. NOTE: This fu
# get_score(importance_type='weight') -> get the feature importance
# -----
# video link2: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

alpha=[10,50,100,500,1000,2000,3000]
cv_log_error_array=[]
for i in alpha:
    x_cfl=XGBClassifier(n_estimators=i)
    x_cfl.fit(X_train_merge,y_train_merge)
    sig_clf = CalibratedClassifierCV(x_cfl, method="sigmoid")
    sig_clf.fit(X_train_merge, y_train_merge)
    predict_y = sig_clf.predict_proba(X_cv_merge)
    cv_log_error_array.append(log_loss(y_cv_merge, predict_y, labels=x_cfl.classes_, e

for i in range(len(cv_log_error_array)):
    print ('log_loss for c = ',alpha[i],'is',cv_log_error_array[i])

best_alpha = np.argmin(cv_log_error_array)

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

x_cfl=XGBClassifier(n_estimators=3000,nthread=-1)
x_cfl.fit(X_train_merge,y_train_merge,verbose=True)
sig_clf = CalibratedClassifierCV(x_cfl, method="sigmoid")
sig_clf.fit(X_train_merge, y_train_merge)

predict_y = sig_clf.predict_proba(X_train_merge)
print ('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_
predict_y = sig_clf.predict_proba(X_cv_merge)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss

```

```

predict_y = sig_clf.predict_proba(X_test_merge)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss)

log_loss for c = 10 is 0.0898979446265
log_loss for c = 50 is 0.0536946658041
log_loss for c = 100 is 0.0387968186177
log_loss for c = 500 is 0.0347960327293
log_loss for c = 1000 is 0.0334668083237
log_loss for c = 2000 is 0.0316569078846
log_loss for c = 3000 is 0.0315972694477

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

For values of best alpha = 3000 The train log loss is: 0.0111918809342
For values of best alpha = 3000 The cross validation log loss is: 0.0315972694477
For values of best alpha = 3000 The test log loss is: 0.0323978515915

```

#### 4.5.5. XgBoost Classifier on final features with best hyper parameters using Random search

```
In [0]: x_cfl=XGBClassifier()
```

```

prams={
    'learning_rate':[0.01,0.03,0.05,0.1,0.15,0.2],
    'n_estimators':[100,200,500,1000,2000],
    'max_depth':[3,5,10],
    'colsample_bytree':[0.1,0.3,0.5,1],
    'subsample':[0.1,0.3,0.5,1]
}
random_cfl=RandomizedSearchCV(x_cfl,param_distributions=prams,verbose=10,n_jobs=-1,)
random_cfl.fit(X_train_merge, y_train_merge)

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

```

[Parallel(n_jobs=-1)]: Done    2 tasks      | elapsed:   1.1min
[Parallel(n_jobs=-1)]: Done    9 tasks      | elapsed:   2.2min
[Parallel(n_jobs=-1)]: Done   19 out of  30 | elapsed:   4.5min remaining:  2.6min
[Parallel(n_jobs=-1)]: Done   23 out of  30 | elapsed:   5.8min remaining:  1.8min
[Parallel(n_jobs=-1)]: Done   27 out of  30 | elapsed:   6.7min remaining:   44.5s
[Parallel(n_jobs=-1)]: Done   30 out of  30 | elapsed:   7.4min finished

```

```
Out[0]: RandomizedSearchCV(cv=None, error_score='raise',
    estimator=XGBClassifier(base_score=0.5, colsample_bylevel=1, colsample_bytree=1,
    gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=3,
    min_child_weight=1, missing=None, n_estimators=100, nthread=-1,
    objective='binary:logistic', reg_alpha=0, reg_lambda=1,
    scale_pos_weight=1, seed=0, silent=True, subsample=1),
    fit_params=None, iid=True, n_iter=10, n_jobs=-1,
    param_distributions={'learning_rate': [0.01, 0.03, 0.05, 0.1, 0.15, 0.2], 'n_
    pre_dispatch='2*n_jobs', random_state=None, refit=True,
    return_train_score=True, scoring=None, verbose=10)
```

```
In [0]: print (random_cfl.best_params_)
```

```
{'subsample': 1, 'n_estimators': 1000, 'max_depth': 10, 'learning_rate': 0.15, 'colsample_bytree': 1}
```

```
In [0]: # find more about XGBClassifier function here http://xgboost.readthedocs.io/en/latest/
# -----
# default paramters
# class xgboost.XGBClassifier(max_depth=3, learning_rate=0.1, n_estimators=100, silent=True,
# objective='binary:logistic', booster='gbtree', n_jobs=1, nthread=None, gamma=0, min_
# max_delta_step=0, subsample=1, colsample_bytree=1, colsample_bylevel=1, reg_alpha=0,
# scale_pos_weight=1, base_score=0.5, random_state=0, seed=None, missing=None, **kwargs)

# some of methods of RandomForestRegressor()
# fit(X, y, sample_weight=None, eval_set=None, eval_metric=None, early_stopping_rounds=None)
# get_params([deep]) Get parameters for this estimator.
# predict(data, output_margin=False, ntree_limit=0) : Predict with data. NOTE: This function
# get_score(importance_type='weight') -> get the feature importance
# -----
# video link2: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

x_cfl=XGBClassifier(n_estimators=1000,max_depth=10,learning_rate=0.15,colsample_bytree=1)
x_cfl.fit(X_train_merge,y_train_merge,verbose=True)
sig_clf = CalibratedClassifierCV(x_cfl, method="sigmoid")
sig_clf.fit(X_train_merge, y_train_merge)

predict_y = sig_clf.predict_proba(X_train_merge)
print ('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(y_train_merge,
predict_y = sig_clf.predict_proba(X_cv_merge)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss(y_cv_merge,
predict_y = sig_clf.predict_proba(X_test_merge)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(y_test_merge,
plot_confusion_matrix(y_test_asm,sig_clf.predict(X_test_merge))
```

```
For values of best alpha = 3000 The train log loss is: 0.0121922832297
```

```
For values of best alpha = 3000 The cross validation log loss is: 0.0344955487471
```

For values of best alpha = 3000 The test log loss is: 0.0317041132442

---

### 1.1.1 With Image Features and 2,3,4 -grams

```
In [0]: # final_data = pd.read_csv('final_data_img.csv')
        final_data = pd.read_csv('final_data_img4.csv')

In [0]: X_train_final, X_test_final, y_train_final, y_test_final = train_test_split(final_data
        X_trn_final, X_cv_final, y_trn_final, y_cv_final = train_test_split(X_train_final, y_t

In [0]: x_cfl = XGBClassifier()

        prams={
            'learning_rate':[0.01,0.03,0.05,0.1,0.15,0.2],
            'n_estimators':[100,200,500,1000,2000],
            'max_depth':[3,5,10],
            'colsample_bytree':[0.1,0.3,0.5,1],
            'subsample':[0.1,0.3,0.5,1]
        }
        random_cfl=RandomizedSearchCV(x_cfl,param_distributions=prams,verbose=10,n_jobs=-1,)
        random_cfl.fit(X_trn_final, y_trn_final)
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 24 concurrent workers.
[Parallel(n_jobs=-1)]: Done   3 out of  30 | elapsed:   55.2s remaining:   8.3min
[Parallel(n_jobs=-1)]: Done   7 out of  30 | elapsed:   2.8min remaining:   9.1min
[Parallel(n_jobs=-1)]: Done  11 out of  30 | elapsed:   3.0min remaining:   5.2min
[Parallel(n_jobs=-1)]: Done  15 out of  30 | elapsed:   3.0min remaining:   3.0min
[Parallel(n_jobs=-1)]: Done  19 out of  30 | elapsed:   8.7min remaining:   5.0min
[Parallel(n_jobs=-1)]: Done  23 out of  30 | elapsed:  11.7min remaining:   3.6min
[Parallel(n_jobs=-1)]: Done  27 out of  30 | elapsed:  13.2min remaining:   1.5min
[Parallel(n_jobs=-1)]: Done  30 out of  30 | elapsed:  23.5min finished
```

```
Out[0]: RandomizedSearchCV(cv='warn', error_score='raise-deprecating',
        estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=
        colsample_bynode=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
        max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
        n_estimators=100, n_jobs=1, nthread=None,
        objective='binary:logistic', random_state=0, reg_alpha=0,
        reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
        subsample=1, verbosity=1),
        fit_params=None, iid='warn', n_iter=10, n_jobs=-1,
        param_distributions={'learning_rate': [0.01, 0.03, 0.05, 0.1, 0.15, 0.2], 'n
        pre_dispatch='2*n_jobs', random_state=None, refit=True,
        return_train_score='warn', scoring=None, verbose=10)
```

```
In [0]: random_cfl.best_params_
```

```
Out[0]: {'subsample': 0.5,  
        'n_estimators': 500,  
        'max_depth': 5,  
        'learning_rate': 0.03,  
        'colsample_bytree': 0.5}
```

```
In [4]: x_cfl=XGBClassifier(n_estimators=500,max_depth=5,learning_rate=0.03,colsample_bytree=0  
x_cfl.fit(X_trn_final,y_trn_final,verbose=True)  
sig_clf = CalibratedClassifierCV(X_cfl, method="sigmoid")  
sig_clf.fit(X_trn_final, y_trn_final)
```

```
# filename = 'finalized_model_clf.sav'  
# pickle.dump(sig_clf, open(filename, 'wb'))  
predict_y = sig_clf.predict_proba(X_trn_final)  
# print ('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log  
print ("The train log loss is:", log_loss(y_trn_final, predict_y))
```

```
predict_y = sig_clf.predict_proba(X_cv_final)  
# print('For values of best alpha = ', alpha[best_alpha], "The cross validation log lo  
print("The cross validation log loss is:", log_loss(y_cv_final, predict_y))
```

```
predict_y = sig_clf.predict_proba(X_test_final)  
# print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log  
print("The test log loss is:", log_loss(y_test_final, predict_y))
```

```
# plot_confusion_matrix(y_test_asm,sig_clf.predict(x_test_final))
```

The train log loss is: 0.009331682518494231

The cross validation log loss is: 0.01211256309293162

The test log loss is: 0.01623229063638662

```
In [5]: from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["Algorithm", "Train log loss", "Test log loss"]
```

```
x.add_row(["Random", 2.45615644965, 2.48503905509])  
x.add_row(["KNN", 0.0782947669247, 0.241508604195 ])  
x.add_row(["Logistic Regression", 0.498923428696,0.528347316704 ])  
x.add_row(["Random Forest", 0.0166267614753, 0.0401141303589 ])  
x.add_row(["XGBOOST", 0.0111918809342, 0.0323978515915 ])  
x.add_row(["XGBOOST-Random Search", 0.0121922832297, 0.0317041132442 ])  
x.add_row(["XGBOOST-(Image features/N-Grams)", 0.009331682518494231,0.01623229063638662])
```

```
print(x)
```

Algorithm	Train log loss	Test log loss
Random	2.45615644965	2.48503905509
KNN	0.0782947669247	0.241508604195
Logistic Regression	0.498923428696	0.528347316704
Random Forest	0.0166267614753	0.0401141303589
XGBOOST	0.0111918809342	0.0323978515915
XGBOOST-Random Search	0.0121922832297	0.0317041132442
XGBOOST-(Image features/N-Grams)	0.009331682518494231	0.01623229063638662